

Statistical Methods

Samy Braik

We denote by p the target distribution and q an easy-to-sample distribution, for example a centered Gaussian.


Let $X_0 \sim p$. We want to add noise until we reach pure noise, and denoise it afterward. We choose an horizon of time $T \in \mathbb{N}^*$ and a noise schedule $\beta : [0, T] \rightarrow \mathbb{R}^*$, continuous and non decreasing.

Forward process

$$d\vec{X}_t = \frac{-\beta(t)}{2\sigma^2} \vec{X}_t dt + \sqrt{\beta(t)} dB_t, \quad \vec{X}_0 \sim p \quad (1)$$

Backward process

$$\begin{aligned} d\overleftarrow{X}_t = & \left(\frac{\beta(T-t)}{2\sigma^2} \overleftarrow{X}_t + \beta(T-t) \nabla \log p_{T-t} \left(\overleftarrow{X}_t \right) \right) dt \\ & + \sqrt{\beta(T-t)} dB_t, \quad \overleftarrow{X}_0 \sim p_T \end{aligned} \quad (2)$$



We learn the score by using score-matching techniques

Score matching

$$\mathcal{L}_{\text{score}}(\theta) = \mathbb{E} \left[\left\| s_{\theta} \left(\tau, \vec{X}_{\tau} \right) - \log p_{\tau} \left(\vec{X}_{\tau} | X_0 \right) \right\|^2 \right] \quad (3)$$

where s_{θ} is a neural network approximating the score function.

Plug it in the backward process and generate by discretizing the dynamics.

Let $X_0 \sim q$ and $f : \mathbb{R}^d \rightarrow \mathbb{R}^d$ an invertible and differentiable function an set $X_1 := f(X_0)$ such that $X_1 \sim p$.

We can write the density of X_1 as

$$p_{X_1} = p_{X_0}(f^{-1}(x_1)) \left| \det \frac{\partial f^{-1}}{\partial x_1}(x_1) \right| \quad (4)$$

$$= p_{X_0}(f^{-1}(x_1)) \left| \det \frac{\partial f}{\partial x_0}(f^{-1}(x_1)) \right|^{-1} \quad (5)$$

We can then write the log-likelihood as

$$\log p_{X_1}(x_1) = \log p_{X_0}(f^{-1}(x_1)) - \log \left| \det \frac{\partial f}{\partial x_0}(f^{-1}(x_1)) \right| \quad (6)$$

Let $X_0 \sim q$ and $X_1 \sim p$. We want to learn f_θ such that $X_1 \simeq f_\theta(X_0) = Z \sim p_Z$. To do that, we set a structure on f_θ , with f_1, \dots, f_k simpler function (all parametrized by θ) such that

$$f_\theta = f_1 \circ f_2 \circ \dots \circ f_k$$

We determine f_θ by minimizing

$$\mathcal{L}_{\text{NF}}(\theta) = \mathbb{E} \left[-\log p_Z(f_\theta(x)) - \log \left| \det \frac{\partial f_\theta}{\partial x}(x) \right| \right]$$



We start by defining a probability density path



Probability density path



A probability path $p : [0, 1] \times \mathbb{R}^d \rightarrow \mathbb{R}^d$ meaning that for each time t , p_t is density function i.e. $\int p_t(x) dx = 1$.

A simple example of such a path is a path p interpolating two density p_0 and p_1 with $p_t = tp_1 + (1 - t)p_0$

Figure: A probability path interpolating $\mathcal{N}(0, 0.2)$ and $\mathcal{U}([0, 1])$







Next we introduce a core object, a time dependent vector field $v : [0, 1] \times \mathbb{R}^d \rightarrow \mathbb{R}^d$ which is used to construct a map $\phi : [0, 1] \times \mathbb{R}^d \rightarrow \mathbb{R}^d$, called a flow, by the following ODE

$$\begin{aligned}\frac{d}{dt}\phi_t(x) &= v_t(\phi_t(x)) \\ \phi_0(x) &= x\end{aligned}\tag{7}$$

The link between the flow and the probability path is given by the change of variable formula

$$p_t(x) = q(\phi_t^{-1}(x)) \det \left[\frac{\partial \phi_t^{-1}}{\partial x}(x) \right]\tag{8}$$

That coincides with the normalizing flow case.

The link between the vector field and the probability path is given by the continuity equation

$$\frac{d}{dt}p_t(x) + \operatorname{div}[v_t(x)p_t(x)] = 0 \quad (9)$$

It said that the vector field v_t generates the probability path p_t if it satisfies the continuity equation.




Figure: Link between marginal flow matching objects

Given a probability path p_t and a vector field v_t , the naïve flow matching loss is defined by

$$\mathcal{L}_{\text{FM}}(\theta) = \mathbb{E}_{t, p_t(x)} \left[\|v_t^\theta(x)t - v_t(x)\|^2 \right] \quad (10)$$

but we don't have access to v_t and p_t . To address this problem, we introduce new objects.

Given a particular data sample x_1 from q , we introduce a conditional probability path $p_t(x|x_1)$ such that at time $t = 0$ $p_0(x|x_1) = q(x)$ and by marginalizing over x_1 we can recover the marginal probability path

$$p_t(x) = \int p_t(x|x_1)q(x_1)dx_1 \quad (11)$$

Instead of working with a probability density, we consider samples from our distributions and define how to go from one to the other.

In the same vein, we can define a conditional vector field, assuming $p_t(x) > 0$ for all t and x

$$v_t(x) = \int v_t(x|x_1) \frac{p_t(x|x_1)q(x_1)}{p_t(x)} dx_1 \quad (12)$$




Figure: Link between all the flow matching objects

The new loss function, called conditional flow matching loss, is defined as

$$\mathcal{L}_{\text{CFM}}(\theta) = \mathbb{E}_{t, p_t(x|x_1)} \left[\|v_t^\theta(x|x_1)t - v_t(x|x_1)\|^2 \right] \quad (13)$$

with the property : $\mathcal{L}_{\text{FM}}(\theta) = \mathcal{L}_{\text{CFM}}(\theta)$ up to a constant independant of θ .

Algorithm Flow matching training

Input: dataset p , noise q

Initialized v^θ

while not converged **do**

$t \sim \mathcal{U}([0, 1])$

$x_1 \sim p(x_1)$

$x_0 \sim q(x_0)$

 Gradient step with $\nabla_\theta \|v_t^\theta(x_t - \dot{x}_t)^2\|$

Output: v^θ

Algorithm Flow matching sampling

Input: Trained v^θ

$x_0 \sim q(x_0)$

Solve numerically the ODE $\dot{x}_t = v_t^\theta(x_t)$

Output: x_1

Comparison

Models	Pros	Cons
Diffusion	1.2	Only works with Gaussian
Normalizing flow	Exact density estimation	Computationally intensive
Flow matching	Exact density estimation	test
	Simulation free training	Slow rate of convergence
Kernel estimator	Flexible Easy to exploit	Hard to evaluate at new data point
		Hard to choose tuning parameters
		Need a lot of data